

Mean rating per criterion (1-5)

	Faithfulness		Both	Technical quality					Semantic quality			
	grounding	precision	numerical fidelity	linguistic	search bar realism	specificity	phrasing economy	focus	search realism	lexical distance	conceptual framing	retrievability
GPT-5.4-mini	4.83	4.40	4.97	4.57	3.67	3.56	4.72	4.94	4.58	4.33	4.25	4.42
GPT-5-mini	4.80	4.50	4.93	4.60	4.10	4.05	4.75	4.90	4.50	4.10	4.80	4.60
Gemini 3.5 Flash	4.37	3.90	4.83	4.70	4.53	4.65	4.76	4.88	4.38	4.77	3.85	3.85
Qwen3.6-35B-A3B	4.27	4.00	4.97	4.60	4.44	4.56	4.94	4.94	4.57	4.71	3.71	3.29
Grok 4.3	4.50	4.07	5.00	4.43	4.12	4.31	4.94	4.94	4.43	4.43	4.14	3.93
Sonnet 4.6	4.30	3.93	4.97	4.53	4.67	4.87	4.80	5.00	4.93	4.87	4.33	3.67

Faithfulness & linguistic pools: all 30 questions per model · technical / semantic quality pools: questions scored with that rubric (56-64 / 36-56 ratings per cell)
Human annotation, single annotator, 30 patent documents × 6 models (180 questions) · export 2026-08-18