


Overall average score

composite = mean of faithfulness, quality and linguistic sub-scores (1-5 scale, axis zoomed to 4.2-4.8)



Confidence intervals overlap heavily — average-score differences between the models are not statistically meaningful at $n=30$
Human annotation, single annotator, 30 patent documents $\times$ 6 models (180 questions) · export 2026-08-18